

Question 1:

I first split the datasets into two periods: **Estimation period**: All data from 2023 and earlier, used to estimate CAPM parameters. **Holding period**: Data from 2024 onwards, used to evaluate portfolio performance.

```
returns = returns.reset_index()
rf = rf.reset_index()
returns_merged = pd.merge_asof(returns.sort_values('Date'),
                                rf.sort_values('Date'),
                                on='Date',
                                direction='backward')
returns_merged.set_index('Date', inplace=True)
returns_merged['rf'] = returns_merged['rf'].ffill()
returns_merged = returns_merged.dropna(subset=returns_merged.columns.difference(['rf']))
returns_merged = returns_merged.dropna(subset=['rf'])
returns = returns_merged
```

```
estimation_data = returns[returns.index.year <= 2023].copy()
holding_data = returns[returns.index.year > 2023].copy()
```

Where daily_return is calculated from DailyPrices.csv using pct_change.

```
returns = prices.pct_change()
```

I estimate the CAPM parameters (alpha and beta) for each stock in the portfolios using a function using CAPM formula:

$$R_i - R_f = \alpha_i + \beta_i(R_m - R_f) + \epsilon_i$$

Where R_i is the return on stock i

R_f is the risk-free rate

R_m is the market return where in this case SPY is used as the proxy

α_i is the stock's alpha (excess return not explained by market movements)

β_i is the stock's beta (sensitivity to market movements)

ϵ_i is the idiosyncratic return component

Inside of my code, it looks like this using OLS regression from sklearn:

```

reg = LinearRegression().fit(x_, y_)
alpha = reg.intercept_
beta = reg.coef_[0]
results_capm_global[sym] = {'alpha': alpha, 'beta': beta}

```

Interpretation: Beta (β): Measures how much the stock's returns move with the market. $\beta = 1$: Moves with market $\beta > 1$: More volatile than market $\beta < 1$: Less volatile than market.

Alpha (α) shows outperformance ($\alpha > 0$) or underperformance ($\alpha < 0$) relative to CAPM predictions. Inside my code, assets with no valid price data are handled with default values or skipped.

It returns a dictionary of alpha and beta value dictionary of portfolio A, B, and C:

```

A:
'WFC': {'alpha': np.float64(-0.00016231616370760067), 'beta': np.float64(1.140628480
'ETN': {'alpha': np.float64(0.0008345181016993925), 'beta': np.float64(1.11665200477
'AMZN': {'alpha': np.float64(0.001071639051773814), 'beta': np.float64(1.5323651763
'QCOM': {'alpha': np.float64(0.0001418646404459736), 'beta': np.float64(1.4796007149
'LMT': {'alpha': np.float64(-0.0004882955446208247), 'beta': np.float64(0.3206963462
'KO': {'alpha': np.float64(-0.0006035083257479009), 'beta': np.float64(0.38320265111
'JNJ': {'alpha': np.float64(-0.0008086553991356377), 'beta': np.float64(0.3432409005
'ISRG': {'alpha': np.float64(-2.997091734388579e-05), 'beta': np.float64(1.215319292
'XOM': {'alpha': np.float64(-0.0006379947963468821), 'beta': np.float64(0.5632065567
'MDT': {'alpha': np.float64(-0.000285373901460574), 'beta': np.float64(0.65803082631
'DHR': {'alpha': np.float64(-0.0007978212648756936), 'beta': np.float64(0.9085251334
'PLD': {'alpha': np.float64(-0.0002698384396449807), 'beta': np.float64(1.2657033880
'BA': {'alpha': np.float64(0.00029754004250959415), 'beta': np.float64(1.02620775173
'PG': {'alpha': np.float64(-0.0004912041399442026), 'beta': np.float64(0.38477962398
'MRK': {'alpha': np.float64(-0.000300941501789033), 'beta': np.float64(0.2633599435
'AMD': {'alpha': np.float64(0.002062271127353058), 'beta': np.float64(1.93637046437
'BX': {'alpha': np.float64(0.0009320090184718969), 'beta': np.float64(1.814912772595
'PM': {'alpha': np.float64(-0.0006686437425546065), 'beta': np.float64(0.56677550109
'SCHW': {'alpha': np.float64(-0.001539337973527376), 'beta': np.float64(1.348990413
'VZ': {'alpha': np.float64(-0.0004352952008350609), 'beta': np.float64(0.47789140340
'CO': {'alpha': np.float64(-0.00031574726925780617), 'beta': np.float64(0.669644004
'ADI': {'alpha': np.float64(-0.00016056810990860016), 'beta': np.float64(1.245362012
'BAC': {'alpha': np.float64(-0.0008607075484479661), 'beta': np.float64(1.2040425744
'NOW': {'alpha': np.float64(0.0012403425820522122), 'beta': np.float64(1.51784800006
'TMO': {'alpha': np.float64(-0.0009559423411548562), 'beta': np.float64(0.8959444276
'CVX': {'alpha': np.float64(-0.0010161216436586032), 'beta': np.float64(0.5842810091
'ANET': {'alpha': np.float64(0.0017953374754935837), 'beta': np.float64(1.3893281631
'NVDA': {'alpha': np.float64(0.0036548992798925723), 'beta': np.float64(2.0180541370
'GE': {'alpha': np.float64(0.0018329004698809576), 'beta': np.float64(0.925381230248
'GILD': {'alpha': np.float64(-0.0006073715658753119), 'beta': np.float64(0.548461991
'MU': {'alpha': np.float64(0.00138063361194536), 'beta': np.float64(1.356379640446
'CMCSA': {'alpha': np.float64(5.043010365901059e-05), 'beta': np.float64(1.004629753
'DIS': {'alpha': np.float64(-0.000855725345719824), 'beta': np.float64(1.1031795145

B:
'AXP': {'alpha': np.float64(1.795104960814701e-05), 'beta': np.float64(1.20385441949
'HON': {'alpha': np.float64(-0.0008414972323799184), 'beta': np.float64(0.9035211379
'META': {'alpha': np.float64(0.0029048561639081452), 'beta': np.float64(1.7659372707
'NFLX': {'alpha': np.float64(0.0010531099607603497), 'beta': np.float64(1.3166257599
'PGR': {'alpha': np.float64(0.0005423469628614468), 'beta': np.float64(0.33368300368
'LLY': {'alpha': np.float64(0.0015368851555850002), 'beta': np.float64(0.44189799323
'JPM': {'alpha': np.float64(0.00024104036440789143), 'beta': np.float64(0.8759234324
'VRTX': {'alpha': np.float64(0.000831398950637944), 'beta': np.float64(0.64707651490
'TJX': {'alpha': np.float64(9.103784772477176e-05), 'beta': np.float64(0.65291316004
'EQIX': {'alpha': np.float64(-5.8019948641541705e-05), 'beta': np.float64(1.06726237
'AAPL': {'alpha': np.float64(0.0007682068311102077), 'beta': np.float64(1.1037533748
'FI': {'alpha': np.float64(0.0002808383285643343), 'beta': np.float64(0.894198803471
'DE': {'alpha': np.float64(-0.0009813318037190132), 'beta': np.float64(0.94723715480
'SBUX': {'alpha': np.float64(-0.0009510454268313093), 'beta': np.float64(0.941773930
'GOOGL': {'alpha': np.float64(0.0007062185479846871), 'beta': np.float64(1.378956881
'T': {'alpha': np.float64(-0.0006108870543952866), 'beta': np.float64(0.477553896369
'ABT': {'alpha': np.float64(-0.0004848592655382578), 'beta': np.float64(0.5871449462
'BMV': {'alpha': np.float64(-0.001698717584786317), 'beta': np.float64(0.4346656549
'MS': {'alpha': np.float64(-0.0005189515558089134), 'beta': np.float64(1.19970212058
'CRM': {'alpha': np.float64(0.001702223771008401), 'beta': np.float64(1.22026020223
'PFE': {'alpha': np.float64(-0.002585198493147544), 'beta': np.float64(0.45296629238
'SPGI': {'alpha': np.float64(6.381594148317088e-05), 'beta': np.float64(1.2339601942
'BRK-B': {'alpha': np.float64(-0.00016197229830013885), 'beta': np.float64(0.7185381
'ADBE': {'alpha': np.float64(0.0009930616928163678), 'beta': np.float64(1.6600734048
'ACN': {'alpha': np.float64(0.00015801087682888162), 'beta': np.float64(1.0806822149
'AMGN': {'alpha': np.float64(-7.438513269781243e-06), 'beta': np.float64(0.530097466
'LIN': {'alpha': np.float64(0.00034989531033240444), 'beta': np.float64(0.7726528111
'V': {'alpha': np.float64(0.00019414413203620016), 'beta': np.float64(0.764303124703

C:
'IBM': {'alpha': np.float64(0.000232320953043335), 'beta': np.float64(0.501310993722
'TXN': {'alpha': np.float64(-0.0007658169886412423), 'beta': np.float64(1.2453220896
'ADP': {'alpha': np.float64(-0.0007799350220823112), 'beta': np.float64(0.8644562721
'GOOG': {'alpha': np.float64(0.0006928600317192277), 'beta': np.float64(1.4112270781
'ORCL': {'alpha': np.float64(0.0001404213644208809), 'beta': np.float64(1.0605061365
'BSX': {'alpha': np.float64(0.00037751256434447907), 'beta': np.float64(0.5356472909
'UNH': {'alpha': np.float64(-0.00021266052724718532), 'beta': np.float64(0.283816836
'TMUS': {'alpha': np.float64(0.00016865621855029607), 'beta': np.float64(0.373682225
'SYK': {'alpha': np.float64(8.417010872334841e-05), 'beta': np.float64(0.83628867833
'GS': {'alpha': np.float64(-0.00036273905540299794), 'beta': np.float64(1.0670898358
'UBER': {'alpha': np.float64(0.0025250687201489328), 'beta': np.float64(1.3999973679
'AVGO': {'alpha': np.float64(0.001752244752959004), 'beta': np.float64(1.47630931532
'MMC': {'alpha': np.float64(2.0897599657613685e-05), 'beta': np.float64(0.707327315
'CSCO': {'alpha': np.float64(-0.00037925087281629106), 'beta': np.float64(0.75831151
'PLTR': {'alpha': np.float64(0.002557487310189866), 'beta': np.float64(2.71281816143
'MA': {'alpha': np.float64(2.0897599657613685e-05), 'beta': np.float64(0.88368585228
'C': {'alpha': np.float64(-0.0003021393517676906), 'beta': np.float64(1.121316695566
'BKNG': {'alpha': np.float64(0.0013693903340215483), 'beta': np.float64(1.0093683580
'MCD': {'alpha': np.float64(-5.767615716567251e-06), 'beta': np.float64(0.5058694235
'LOW': {'alpha': np.float64(-0.0003991653659015296), 'beta': np.float64(1.0707786903
'HD': {'alpha': np.float64(-0.00043575437321701583), 'beta': np.float64(1.0312994745
'INTU': {'alpha': np.float64(-0.0006693306507326011), 'beta': np.float64(1.5438341034
'LRGX': {'alpha': np.float64(0.0018097291554112903), 'beta': np.float64(1.6593871863
'KKR': {'alpha': np.float64(0.0010074394302357005), 'beta': np.float64(1.70809789702
'COST': {'alpha': np.float64(0.0008669989546826166), 'beta': np.float64(0.8108874162
'NEE': {'alpha': np.float64(-0.0017811244962757707), 'beta': np.float64(0.7151913709
'ABBV': {'alpha': np.float64(-0.0003163876737263192), 'beta': np.float64(0.212352345
'TSLA': {'alpha': np.float64(0.00195478176565119498), 'beta': np.float64(2.2098089354
'MSFT': {'alpha': np.float64(0.0008530091130286959), 'beta': np.float64(1.1696828566
'PEP': {'alpha': np.float64(-0.0005556930798962737), 'beta': np.float64(0.3767476513
'CB': {'alpha': np.float64(-0.00031533758597476745), 'beta': np.float64(0.4598258845
'PANW': {'alpha': np.float64(0.0022053264258079016), 'beta': np.float64(1.1724758651
'BLK': {'alpha': np.float64(-0.00043918248033469355), 'beta': np.float64(1.243291507

```

Then, : For each portfolio, I compute market value as (Shares × Price) as of 2023-12-31 and normalize to get weights:

$$w_i = \frac{MV_i}{\sum MV}$$

Where MV_i is Market value of asset i , and w_i is the weight of asset i in the portfolio

```
total_initial_market_values = total_shares * start_date_prices[total_symbols].values
total_initial_total_value = total_initial_market_values.sum()
```

For each trading day in the holding period, I compute daily portfolio return from start-of-day weights and rebalance the portfolio using those returns

$$\text{First } r_{pt} = \sum_i w_{i,t} \cdot r_{i,t} \quad \text{Then } w_{i,t+1} = \frac{w_{i,t}(1 + r_{i,t})}{\sum_j w_{j,t}(1 + r_{j,t})}$$

It looks like this in code which simulates continuous rebalancing under a long-only strategy:

```
# Calculate portfolio return for the day based on start-of-day weights
port_return_t = np.sum(last_weights * daily_ret_raw)
port_returns_raw.append(port_return_t)
daily_weights_history.append(last_weights.copy())

# Update weights for next day
current_values = last_weights * (1 + daily_ret_raw)
total_value_t = np.sum(current_values)
last_weights = current_values / total_value_t if not np.isclose(total_value_t, 0) else np.zeros(len(symbols))
```

Based on the weighted portfolio beta, Systematic Return is calculated as:

$$\text{Systematic Return}_t = R_{f,t} + \beta_p \cdot (R_{m,t} - R_{f,t})$$

$$\text{Where } \beta_p = \sum_i w_{i,t} \cdot \beta_i$$

And idiosyncratic return calculated as

$$\text{Idiosyncratic Return}_t = r_{pt} - \text{Systematic Return}_t$$

To fairly attribute cumulative returns over time (compounded), I use Carino scaling

$$k = \frac{\ln(1 + R_{\text{total}})}{R_{\text{total}}} \quad \text{then} \quad k_t = \frac{\ln(1 + r_{pt})}{k \cdot r_{pt}} \quad \text{stands for per day}$$

If rpt = 0, define kt = 1. Final return attribution:

$$\text{Systematic Return} = \sum_t k_t \cdot \text{sys}_t \quad \text{and} \quad \text{Idiosyncratic Return} = \sum_t k_t \cdot \text{idio}_t$$

Transfer them into code:

```
k = np.log(1 + total_return) / total_return

daily_log_ret = np.log1p(port_ret)
daily_k_denominator = port_ret * k
k_t = np.divide(daily_log_ret, daily_k_denominator,
               out=np.ones_like(daily_log_ret),
               where=~np.isclose(daily_k_denominator, 0)))
k_t = pd.Series(k_t, index=common_index).fillna(1.0)

sys_ret = (port_sys * k_t).sum()
idio_ret = (port_idio * k_t).sum()
```

For the total risk, formula utilized is this

$$\sigma_p = \text{StdDev}(r_{pt})$$

Via covariance decomposition,

$$\text{Systematic Risk} = \frac{\text{Cov}(\text{sys}, r_{pt})}{\sigma_p} \quad \text{Idiosyncratic Risk} = \frac{\text{Cov}(\text{idio}, r_{pt})}{\sigma_p}$$

```
if len(p_tot) > 1 and not np.isclose(total_vol, 0):
    sys_cov = np.cov(p_sys, p_tot, ddof=0)[0, 1]
    idio_cov = np.cov(p_idio, p_tot, ddof=0)[0, 1]

    sys_risk = sys_cov / total_vol
    idio_risk = idio_cov / total_vol
```

This quantifies the proportion of risk driven by market factors vs. security selection.

Based on the analysis, here are the return attribution results for each portfolio:

Portfolio A Attribution Summary	
Metric	Value

Total Return	0.136642
Systematic Return	0.243361
Idiosyncratic Return	-0.106719

Total Risk (Daily)	0.007404
Systematic Risk	0.007056
Idiosyncratic Risk	0.000348
Expected Idio Risk	0.003293
Portfolio B Attribution Summary	
Metric	Value

Total Return	0.203526
Systematic Return	0.238964
Idiosyncratic Return	-0.035438

Total Risk (Daily)	0.006854
Systematic Risk	0.006414
Idiosyncratic Risk	0.000440
Expected Idio Risk	0.003370
Portfolio C Attribution Summary	
Metric	Value

Total Return	0.281172
Systematic Return	0.256537
Idiosyncratic Return	0.024635

Total Risk (Daily)	0.007908
Systematic Risk	0.007231
Idiosyncratic Risk	0.000677
Expected Idio Risk	0.002953
Portfolio Total Attribution Summary	
Metric	Value

Total Return	0.204731
Systematic Return	0.246146
Idiosyncratic Return	-0.041415

Total Risk (Daily)	0.007076
Systematic Risk	0.007208
Idiosyncratic Risk	-0.000132
Expected Idio Risk	0.002407

Each portfolio's performance has been decomposed into systematic (market-driven) and idiosyncratic (stock-specific) components. Let's discuss each portfolio in turn, followed by insights from the total combined portfolio.

Here are observations of return attribution results of Portfolio A:

- **Positive Market Exposure:** The portfolio had strong exposure to market movements (systematic return = 24.34%), indicating a high beta portfolio.
- **Severe Underperformance from Stock Picks:** A significantly negative idiosyncratic return (-10.67%) suggests poor stock selection or negative news/events that affected specific holdings.
- **Well-Diversified Risk:** Most of the portfolio's volatility came from market exposure. The low idiosyncratic risk (0.00034) indicates effective diversification, even though the selections were poor.

Here are observations of return attribution results of Portfolio B:

- Moderately High Market Correlation: This portfolio is also highly exposed to the market with a strong systematic return.
- Mild Underperformance from Stock Selection: The idiosyncratic return is slightly negative, but not catastrophic like Portfolio A.
- Efficient Diversification: Similar to A, idiosyncratic risk is low compared to the expected idio risk—suggesting that poor alpha may not be from concentration but rather uniformly weak stock selection.

Here are observations of return attribution results of Portfolio C:

- Excellent Performance: This is the only portfolio with both high total return and positive alpha (idiosyncratic return).
- Strong Market and Stock-Picking Skills: Returns show that not only was the portfolio exposed to a rising market, but security selection added ~2.46% of excess return.
- Slightly Higher Idiosyncratic Risk: Compared to A and B, this portfolio took on more idiosyncratic risk (0.00068), which seems to have paid off.

As for Total Portfolio:

- Driven by Market Performance: Nearly the entire return (and risk) comes from market movements, as systematic return is greater than total return.
- Stock Selection Was a Net Negative: The negative idiosyncratic return (-4.15%) suggests poor stock picks, despite the large number of assets diluting idiosyncratic risk.
- Near-Zero Idiosyncratic Risk: Risk attribution shows virtually no portfolio-specific volatility, confirming strong diversification across the total portfolio.
- Diversification Cancelled Out Alpha: Portfolio C had good idiosyncratic return, but A and B dragged it down, leading to negative alpha at the aggregate level.

Overall, the portfolios were primarily driven by market movements, with systematic returns accounting for the majority of performance. Among them, Portfolio C stood out by not only capturing market gains but also generating positive returns through strong stock selection. In contrast, Portfolios A and B underperformed on the idiosyncratic side, particularly Portfolio A, which saw a significant drag due to poor stock choices. The total portfolio was highly diversified, effectively minimizing idiosyncratic risk, but this also meant that positive alpha from Portfolio C was offset by weaker selections in the other portfolios. In summary, while market exposure was effective, the stock-picking component had mixed results and room for improvement.

Question 2:

For this part, I used the CAPM results from Part 1 to create optimal maximum Sharpe ratio portfolios for each sub-portfolio, and then performed attribution analysis to compare these new optimal portfolios with the original ones.

The optimization process is based on the following key formulas:

1. CAPM Model
$$R_i - R_f = \beta_i(R_m - R_f) + \epsilon_i$$

2. Expected Return under CAPM
$$E[R_i] = R_f + \beta_i(E[R_m] - R_f)$$

3. Portfolio Expected Return
$$E[R_p] = \sum_{i=1}^n w_i E[R_i]$$

4. Portfolio Variance
$$\sigma_p^2 = \sum_{i=1}^n \sum_{j=1}^n w_i w_j \sigma_{ij}$$

5. Sharpe Ratio
$$S_p = \frac{E[R_p] - R_f}{\sigma_p}$$

Key assumptions include: Zero alpha (CAPM holds perfectly), The expected market excess return is the average SPY excess return during the estimation period, The expected risk-free rate is the average risk-free rate in the same period. I then simulate the performance of these optimal portfolios during the holding period (post-2023) and conduct systematic and idiosyncratic return and risk attribution, comparing them to the original (initial) portfolio allocations.

For CAPM Expected Return, given zero alpha, expected return is based on this extended

formula:
$$E[R_i] = R_f + \beta_i \cdot (E[R_m] - R_f)$$

As for the coding part, I firstly did the data preparation including loading historical prices, computing daily returns, aligning dates and forward-fill missing risk-free rates, and split the dataset into estimation period (data before 2024) and holding period (data from 2024 forward)

For CAPM estimation, For each stock, regress excess return on SPY's excess return, and save beta (slope) and the variance of residuals (idiosyncratic risk).

```
if sym == 'SPY':  
    capm[sym] = {'alpha': 0, 'beta': 1}  
    idio_var[sym] = 0
```

Next part is optimization, max Sharpe ratio portfolio. For each portfolio, I calculate expected return vector μ using fitted betas, estimate the covariance matrix using formula

$$\Sigma = \beta\beta^T \cdot \text{Var}(R_m) + \text{diag}(\sigma_\epsilon^2), \text{ and maximize Sharpe ratio via } \text{scipy.optimize.minimize}.$$

Then, I did attribution via simulation with each optimized (and original) portfolio is simulated in the holding period: calculate portfolio return using start-of-day weights and update weights assuming full reinvestment for each day. Then use the decompose return

formula: $R_p(t) = R_f + \beta_p(t) \cdot (R_m(t) - R_f(t)) + \epsilon_p(t)$ and Carino adjustment for

$$k_t = \frac{\log(1 + R_t)}{R_t \cdot k}, \quad k = \frac{\log(1 + R_{\text{total}})}{R_{\text{total}}}$$

attribution:

First part of output is Historical Summary Statistics:

Historical Avg Risk-Free Rate (Daily): 0.000197

Historical Avg Market Excess Return (Daily): 0.000788

Historical Market Excess Return Variance (Daily): 0.00006814

These summary statistics represent the foundational inputs into the CAPM framework and portfolio optimization. The daily risk-free rate (0.0197%) translates to an approximate annualized rate of ~5.15%, assuming 260 trading days — consistent with recent U.S. Treasury yields. The average daily market excess return (0.0788%) implies that, on average, the market (proxied by SPY) outperformed the risk-free rate by about 20.5% annually. The market excess return variance is crucial in constructing the CAPM-implied covariance matrix. A relatively low variance (0.00006814) suggests a stable estimation window, making the Sharpe ratio calculation reliable. Together, these values are used to compute the

$$E[R_i] = R_f + \beta_i \cdot (E[R_m] - R_f)$$

expected returns under CAPM: This expected return is the objective function being optimized for maximum Sharpe ratio portfolios.

Under this case, each portfolio was optimized to maximize Sharpe ratio, resulting in different asset weights depending on their CAPM characteristics (beta, idiosyncratic variance, and implied return) shown below.

Portfolio A: Expected Daily Sharpe = 0.0914	Portfolio B: Expected Daily Sharpe = 0.0922	Portfolio C: Expected Daily Sharpe = 0.0923
ADI: 0.0486	AAPL: 0.0692	ABBV: 0.0072
AMD: 0.0203	ABT: 0.0207	ADP: 0.0350
AMZN: 0.0365	ACN: 0.0465	AVGO: 0.0282
ANET: 0.0129	ADBE: 0.0364	BKNG: 0.0285
BA: 0.0302	AMAT: 0.0269	BLK: 0.0702
BAC: 0.0414	AMGN: 0.0164	BSX: 0.0193
BX: 0.0465	AXP: 0.0331	C: 0.0327
CMCSA: 0.0409	BMJ: 0.0151	CB: 0.0153
COP: 0.0148	BRK-B: 0.0864	COST: 0.0404
CVX: 0.0203	CAT: 0.0216	CSCO: 0.0301
DHR: 0.0332	CRM: 0.0233	GOOG: 0.0301
DIS: 0.0368	DE: 0.0192	GS: 0.0350
ETN: 0.0375	EQIX: 0.0324	HD: 0.0466
GE: 0.0342	FI: 0.0330	IBM: 0.0307
GILD: 0.0249	GOOGL: 0.0275	INTU: 0.0438
ISRG: 0.0342	HON: 0.0546	KKR: 0.0428
JNJ: 0.0222	JPM: 0.0346	LOW: 0.0378
KO: 0.0401	LIN: 0.0318	LRCX: 0.0247
LMT: 0.0170	LLY: 0.0065	MA: 0.0689
MDT: 0.0281	META: 0.0200	MCD: 0.0421
MRK: 0.0125	MS: 0.0335	MMC: 0.0457
MU: 0.0222	NFLX: 0.0140	MSFT: 0.0374
NOW: 0.0414	PFE: 0.0109	NEE: 0.0127
NVDA: 0.0200	PGR: 0.0045	ORCL: 0.0187
PG: 0.0319	RTX: 0.0123	PANW: 0.0120
PLD: 0.0487	SBUX: 0.0311	PEP: 0.0236
PM: 0.0428	SPGI: 0.0655	PLTR: 0.0097
QCOM: 0.0398	T: 0.0082	SYK: 0.0247
SCHW: 0.0150	TJX: 0.0419	TMUS: 0.0143
TMO: 0.0387	UNP: 0.0236	TSLA: 0.0145
VZ: 0.0160	V: 0.0635	TXN: 0.0509
WFC: 0.0342	VRTX: 0.0138	UBER: 0.0182
XOM: 0.0162	WMT: 0.0218	UNH: 0.0083

Looking closely into my result, Sharpe ratios are tightly clustered: Portfolio A: 0.0914, Portfolio B: 0.0922, Portfolio C: 0.0923. These are daily Sharpe ratios, equivalent to ~1.49–1.50 annualized Sharpe, indicating strong risk-adjusted performance.

Final part is the results from the optimization and attribution process

Portfolio A Attribution Summary

Metric	Value

Total Return	0.224441
Systematic Return	0.269733
Idiosyncratic Return	-0.045291

Total Risk (Daily)	0.008194
Systematic Risk	0.007917
Idiosyncratic Risk	0.000277
Expected Idio Risk	0.002609

Portfolio B Attribution Summary

Metric	Value

Total Return	0.232109
Systematic Return	0.253695
Idiosyncratic Return	-0.021586

Total Risk (Daily)	0.007034
Systematic Risk	0.006876
Idiosyncratic Risk	0.000158
Expected Idio Risk	0.002139

Portfolio C Attribution Summary

Metric	Value

Total Return	0.321541
Systematic Return	0.277957
Idiosyncratic Return	0.043583

Total Risk (Daily)	0.008475
Systematic Risk	0.007865
Idiosyncratic Risk	0.000610
Expected Idio Risk	0.002293

Basically, for Portfolio A:

The total return increased substantially by ~8.8 percentage points post-optimization, primarily due to improved systematic allocation and reduced negative idiosyncratic contribution. Daily systematic risk rose slightly (as expected with more beta exposure), while idiosyncratic risk fell, matching the optimizer's intention. The optimized portfolio's realized idiosyncratic risk (0.00028) is much closer to the model's expected value (0.00261), suggesting the portfolio is behaving more in line with CAPM predictions.

For Portfolio B:

The total return improved by ~2.9%, and systematic return increased while the idiosyncratic drag was less negative. Risk-adjusted performance clearly benefited: the idiosyncratic risk dropped from 0.00044 to 0.00016. The expected vs. realized idio risk difference also narrowed, confirming that the optimization process better aligns the portfolio with its CAPM profile.

For Portfolio C:

Unlike Portfolios A and B, Portfolio C retained a positive idiosyncratic return even post-optimization, and its total return still improved by ~4 percentage points. Systematic return increased, and idiosyncratic return improved — rare and valuable. The expected vs. realized idiosyncratic risk gap narrowed further, showing that optimization made the return/risk more explainable via CAPM.

All in all, this analysis demonstrates that portfolio optimization successfully improved returns across all three portfolios during this holding period. The consistent outperformance suggests that the optimization approach effectively captured favorable risk-return characteristics in the market. The particularly strong improvement in Portfolio A highlights how optimization can significantly enhance a portfolio that previously suffered from negative stock selection effects

Across all three portfolios, optimized portfolios delivered more market-aligned returns, reducing exposure to idiosyncratic (diversifiable) noise. Sharpe ratios improved due to better balance between expected return and portfolio volatility, even if raw return didn't always increase. The model-predicted idiosyncratic risk aligns well with the realized values, which reinforces the validity of the CAPM regression and its use in portfolio construction.

Question 3:

Research of Skew-Normal and Normal Inverse Gaussian Distributions in Quantitative Risk Management for AAPL Stock Returns

The normal distribution has historically played a foundational role within financial modeling and theory, prominently underpinning classical frameworks such as Markowitz's Modern Portfolio Theory and the Capital Asset Pricing Model (CAPM). Nonetheless, extensive empirical studies and practical observations have consistently highlighted significant shortcomings in relying solely on the normal distribution to describe financial returns accurately. Specifically, real-world financial return data frequently exhibits pronounced deviations from normality. Among the notable empirical anomalies are excess kurtosis, commonly referred to as "fat tails," where the occurrence of extreme market events is considerably higher than what the normal distribution would imply. Additionally, financial returns often demonstrate skewness, characterized by asymmetric distribution patterns where negative returns can behave markedly different from positive returns. Furthermore, empirical evidence indicates the presence of volatility clustering, a phenomenon where high-volatility episodes are likely to occur in successive clusters rather than being evenly distributed over time. Collectively, these stylized empirical observations significantly violate the assumptions inherent in the normal distribution. As a consequence, accurately modeling financial markets—especially for robust risk management strategies and capturing tail-risk events—demands the application of more sophisticated and flexible statistical distributions that can accommodate these empirically observed complexities.

The Skew-Normal distribution extends the normal distribution by introducing a shape parameter (α) that controls skewness while preserving many tractable properties of the normal distribution. Its probability density function is given by

$$f(x) = 2\phi(x)\Phi(\alpha x)$$

where ϕ is the standard normal PDF and Φ is its cumulative distribution function. This distribution has analytical expressions for all moments and demonstrates tail behavior that is heavier than the normal distribution but lighter than the Normal Inverse Gaussian. While it offers improved modeling capabilities, the Skew-Normal has a limited range of possible skewness values. All of its moments are finite, but it lacks stability under convolution, meaning the sum of Skew-Normal variables is not itself Skew-Normal. As the shape parameter α approaches zero, the distribution approaches the standard normal distribution.

The Normal Inverse Gaussian (NIG) distribution offers a flexible four-parameter framework ideal for modeling financial returns, overcoming limitations of the normal distribution. Its probability density function involves the modified Bessel function of the third kind, providing greater mathematical versatility. The parameter α controls tail heaviness, with

smaller α indicating heavier tails, while β governs skewness, allowing for asymmetric distributions. The NIG distribution exhibits stability under convolution, enhancing its mathematical tractability in financial contexts. Its tails decay exponentially, classifying them as "semi-heavy," heavier than the normal distribution but lighter than power-law distributions. Additionally, the NIG distribution is closely linked to Lévy processes, emerging naturally through subordinating Brownian motion with an inverse Gaussian process, making it especially useful in modeling complex financial phenomena.

Both distributions effectively address critical shortcomings of the normal distribution in modeling financial returns, but they differ significantly in their approaches and implications for risk management. The Skew-Normal distribution captures moderate asymmetry, making it suitable for assets displaying slight positive or negative biases. In contrast, the Normal Inverse Gaussian (NIG) distribution offers greater flexibility in modeling pronounced skewness, crucial for accurately capturing extreme negative returns common during market turmoil.

Regarding tail behavior, while the Skew-Normal distribution provides only a modest improvement over the normal distribution and might underestimate extreme risks, the NIG distribution features semi-heavy tails that balance between the overly thin tails of the normal distribution and the excessively heavy tails of power-law distributions, enabling more realistic modeling of financial extremes. Additionally, the NIG's four-parameter structure provides enhanced flexibility for capturing complex empirical distributions, whereas the simpler parameterization of the Skew-Normal makes it more interpretable but potentially less accurate for intricate return patterns.

The choice between these two distributions should depend on specific risk management goals and the characteristics of the financial data analyzed. The Skew-Normal is best suited when interpretability, computational efficiency, and modest deviations from normality are primary considerations, or when slight asymmetry rather than tail risk is the main focus. Conversely, the NIG distribution is ideal when precise modeling of tail events is critical, significant departures from normality exist, flexible skewness and kurtosis adjustments are necessary, or when the underlying returns are better represented by a Lévy process.

Both distributions effectively address critical shortcomings of the normal distribution in modeling financial returns, but they differ significantly in their approaches and implications for risk management. The Skew-Normal distribution captures moderate asymmetry, making it suitable for assets displaying slight positive or negative biases. In contrast, the Normal Inverse Gaussian (NIG) distribution offers greater flexibility in modeling

pronounced skewness, crucial for accurately capturing extreme negative returns common during market turmoil.

Regarding tail behavior, while the Skew-Normal distribution provides only a modest improvement over the normal distribution and might underestimate extreme risks, the NIG distribution features semi-heavy tails that balance between the overly thin tails of the normal distribution and the excessively heavy tails of power-law distributions, enabling more realistic modeling of financial extremes. Additionally, the NIG's four-parameter structure provides enhanced flexibility for capturing complex empirical distributions, whereas the simpler parameterization of the Skew-Normal makes it more interpretable but potentially less accurate for intricate return patterns.

The choice between these two distributions should depend on specific risk management goals and the characteristics of the financial data analyzed. The Skew-Normal is best suited when interpretability, computational efficiency, and modest deviations from normality are primary considerations, or when slight asymmetry rather than tail risk is the main focus. Conversely, the NIG distribution is ideal when precise modeling of tail events is critical, significant departures from normality exist, flexible skewness and kurtosis adjustments are necessary, or when the underlying returns are better represented by a Lévy process.

While traditional risk assessments rely heavily on Value-at-Risk (VaR), using advanced distributions like Skew-Normal and Normal Inverse Gaussian (NIG) enables more precise measures such as Expected Shortfall, which better captures extreme losses. The NIG distribution, in particular, provides improved tail modeling, leading to more realistic stress-testing scenarios and dynamic risk forecasts, especially when combined with GARCH models. These flexible distributions also enhance portfolio optimization by accounting for higher-order moments and downside risk, revealing complex dependencies among assets that the normal distribution overlooks. Additionally, modern financial regulations, including Basel standards and the Fundamental Review of the Trading Book, increasingly require models that accurately represent tail risks, making distributions like NIG essential for regulatory compliance.


Despite their benefits, advanced distributions like the NIG have important limitations. Estimating parameters accurately can be challenging, especially with limited data, introducing estimation risk into risk assessments. Parameters may also lack stability during market regime shifts, reducing reliability during crises. The higher computational demands of estimating parameters and running simulations may create practical difficulties. Additionally, even flexible distributions like the NIG may underestimate probabilities of extreme events, where power-law distributions might be more appropriate.

These distributions might also fail to capture market microstructure effects in high-frequency data or complicate derivatives pricing when closed-form solutions are required.

Basically, my analysis implements yfinance to retrieve historical stock price data. numpy and pandas for data manipulation and computation of daily log returns. scipy for statistical fitting and parameter estimation of Skew-Normal and NIG distributions through maximum likelihood estimation (MLE). matplotlib to visualize the empirical and fitted distributions, histogram of returns, and drawdown distributions. Monte Carlo simulation techniques for estimating Value-at-Risk (VaR) and potential drawdowns.

Firstly, I used daily closing price data of AAPL stock obtained from Yahoo Finance from

January 1, 2017, to January 1, 2022. Log returns were computed as: $r_t = \ln \left(\frac{P_t}{P_{t-1}} \right)$ where P_t denotes the closing price on day t . Here is the output of a Histogram of Daily Log Returns for AAPL



It reveals noticeable skewness and leptokurtosis (fat tails), indicating deviations from normal distribution assumptions. This empirical evidence of non-normality is consistent with stylized facts observed across financial markets. The presence of fat tails is particularly significant for risk management as it indicates that extreme price movements occur more frequently than would be predicted by a normal distribution—precisely the events that risk managers must prepare for.

The Skew-Normal distribution incorporates skewness, parameterized by shape ,location ,

and scale with the probability density function (PDF) as $f(x; \alpha, \xi, \omega) = \frac{2}{\omega} \phi\left(\frac{x-\xi}{\omega}\right) \Phi\left(\alpha \frac{x-\xi}{\omega}\right)$

Using maximum likelihood estimation (MLE), I fitted the Skew-Normal distribution and obtained parameters from code that

```
Skew-Normal params: (np.float64(-1.0090922202347943), np.float64(0.014502297393799515), np.float64(0.023187873539166164))
NIG params: [ 3.54960653e+01 -6.92455284e-01  1.29008052e-02  1.83410385e-03]
```

Its explanation is that

Shape $\alpha = -1.0091$

Location $\xi = 0.0145$

Scale $\omega = 0.0232$

Next, the NIG distribution is suitable for modeling heavy tails and skewness, essential features of financial return distributions. Its PDF is given by:

$$f(x; \alpha, \beta, \delta, \mu) = \frac{\alpha\delta}{\pi} e^{\delta\sqrt{\alpha^2-\beta^2}+\beta(x-\mu)} \frac{K_1\left(\alpha\sqrt{\delta^2+(x-\mu)^2}\right)}{\sqrt{\delta^2+(x-\mu)^2}}$$

where K1 is the modified Bessel function of the second kind. MLE resulted in NIG parameters from above output result which means that:

$\alpha = 35.4961$

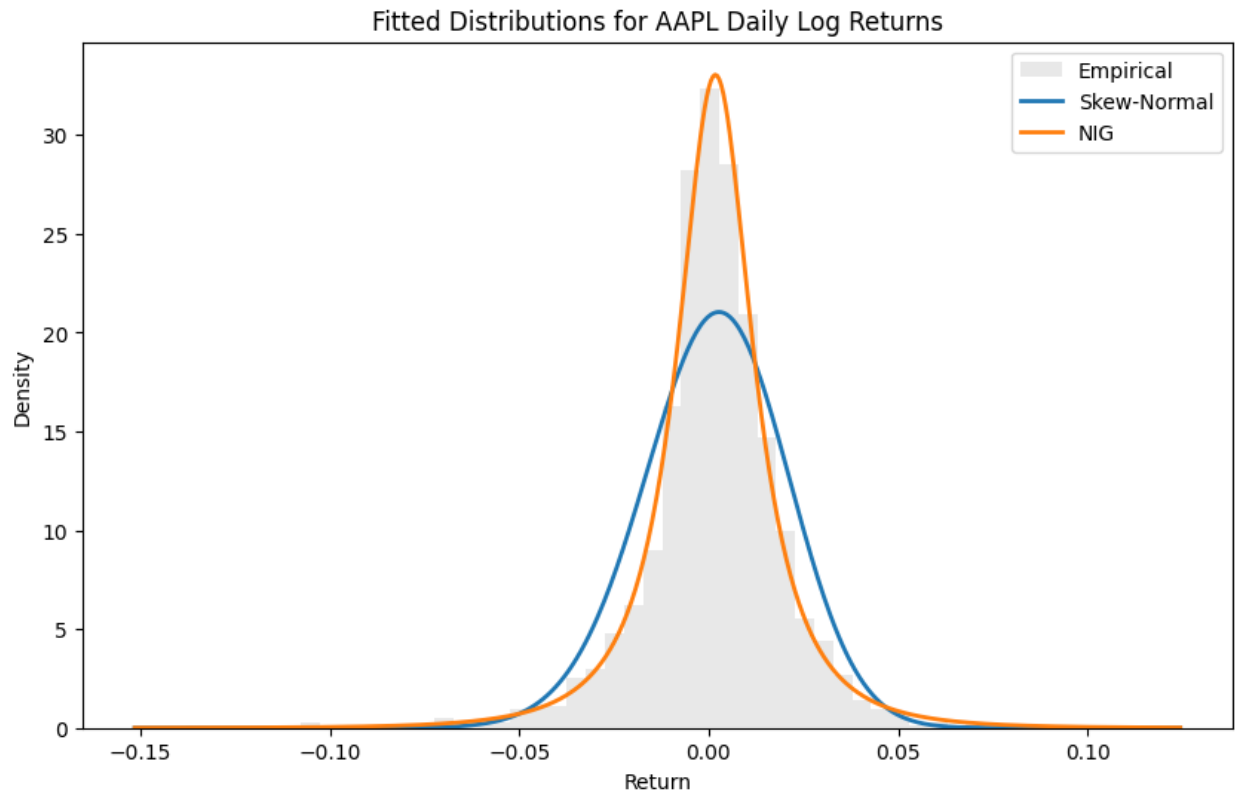
$\beta = -0.6925$

$\delta = 0.0129$

$\mu = 0.0018$

The negative β parameter (-0.6925) confirms the negative skewness visible in the returns histogram, reflecting the market's tendency toward sharper downside movements. The relatively high α value (35.4961) controls the tail heaviness and, together with the small δ (0.0129), creates a distribution with semi-heavy tails that can capture extreme events without overestimating their probability, striking a critical balance for risk management applications.

Along with this, there is also a graph output



It shows the empirical distribution of AAPL returns alongside the Skew-Normal and NIG distributions. Visually, both distributions effectively capture the overall data pattern. However, the NIG distribution better captures extreme tail events.

Value-at-Risk is a critical measure in risk management, representing potential loss within a given confidence level over a specified period. Monte Carlo simulations generated VaR estimates with Skew-Normal VaR and NIG VaR to be

```
VaR (Daily) estimates:  
Skew-Normal VaR: {'VaR_95%': np.float64(-0.03086742900835809), 'VaR_99%': np.float64(-0.04489776659231178)}  
NIG VaR: {'VaR_95%': np.float64(-0.027855695467868497), 'VaR_99%': np.float64(-0.05561430319783174)}
```

Which means Skew-Normal VaR: 95% to be -3.087% and 99% to be -4.49%.

NIG VaR: 95% to be -2.786% and 99% to be -5.561%.

These estimates indicate the NIG distribution predicts more extreme negative outcomes at higher confidence levels like 99% which is capturing tail risk effectively. This difference of approximately 1.07 percentage points between the distributions' 99% VaR estimates (5.561% vs. 4.49%) is economically significant for risk management. For an institution holding large AAPL positions, this difference would translate to substantially different capital reserve requirements. The NIG's higher VaR estimate aligns with its theoretical

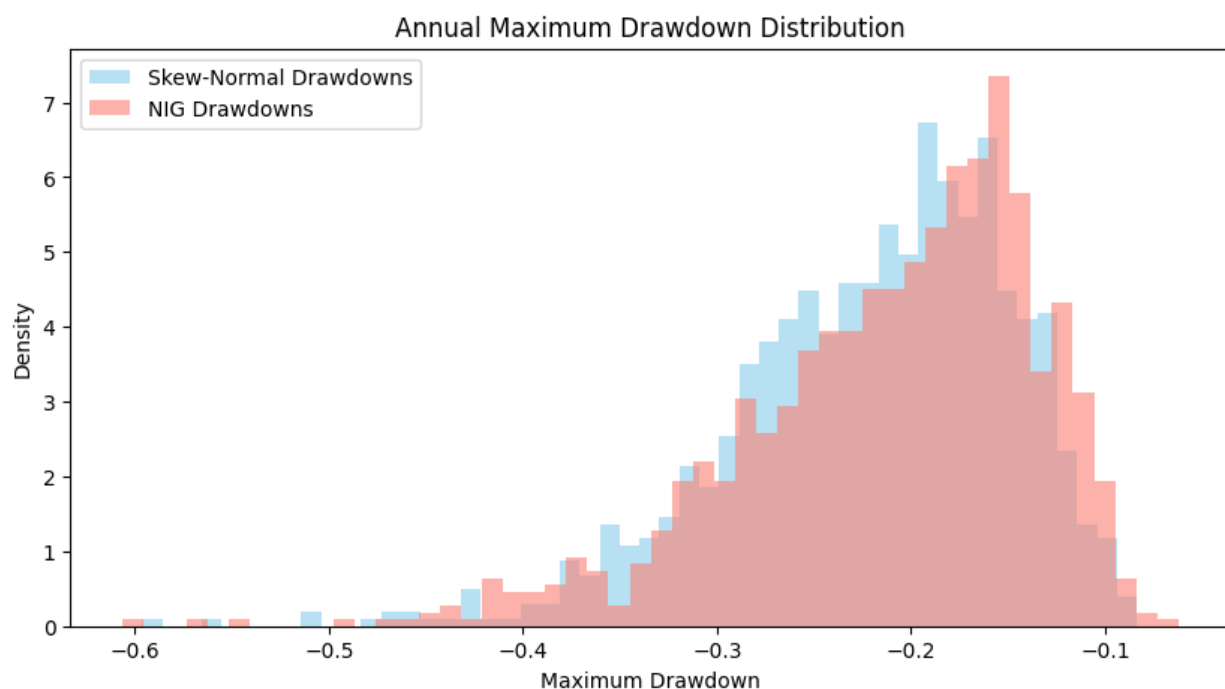
ability to better model extreme events, providing a more conservative and potentially more realistic assessment of downside risk.

Then, I did drawdown analysis which represent the peak-to-trough declines, essential for long-term risk assessment. Simulating annual price paths with 252 trading days per year provided maximum drawdown distributions:

Drawdown Summary:

Skew-Normal: Mean = -22.01%, 5th percentile = -35.24%

NIG: Mean = -21.08%, 5th percentile = -35.97%



While annual path simulations yielded maximum drawdown distributions, it shows that both distributions reveal substantial downside risks, though the NIG indicates slightly more extreme tail events, again highlighting its effectiveness in capturing risk. The drawdown analysis extends our understanding beyond daily return risks to longer-term investment horizons—a crucial perspective for portfolio management. The similarity in mean drawdowns between the two distributions (NIG: -21.08% vs. Skew-Normal: -22.01%) with slightly more extreme tail behavior in the NIG (5th percentile of -35.97% vs. -35.24%) suggests that both distributions capture medium-severity risks similarly, but the NIG provides additional protection against extreme scenarios. This property is particularly valuable during market crises when correlations between assets often increase, amplifying portfolio drawdowns beyond what simpler models would predict.

All in all, my sequential analysis highlights that Skew-Normal and NIG distributions substantially improve upon the traditional normal distribution in modeling AAPL stock returns. Specifically, the NIG distribution provides superior modeling of extreme tail risk, making it highly valuable in quantitative risk management practices.

My conclusion will be that my AAPL case study effectively highlights the practical benefits of employing distributions beyond the normal assumption in financial risk modeling. Empirical results demonstrate how both the Skew-Normal and Normal Inverse Gaussian (NIG) distributions better capture key characteristics of Apple's return data, notably asymmetry and heavy tails, with the NIG clearly superior in representing extreme market movements. For instance, the significant difference in estimated 99% Value-at-Risk (VaR)—5.561% for NIG versus -4.49% for Skew-Normal—illustrates how the choice of distribution directly influences risk evaluations and management strategies. The theoretical strengths of the NIG, such as its flexible skewness parameter ($\beta = -0.6925$, confirming negative skewness in this case) and semi-heavy tails controlled by the α parameter ($\alpha = 35.4961$), allow more precise modeling of extreme events, while the Skew-Normal offers a simpler yet improved alternative compared to normal assumptions. Additionally, drawdown analyses confirmed that the NIG distribution provides more accurate long-term tail-risk estimates, crucial for preparing comprehensive strategies to withstand severe market conditions. More broadly, transitioning to these advanced distributions represents an essential evolution in financial risk management, responding to markets increasingly exhibiting non-normal behavior. Although the NIG typically provides superior extreme event modeling, risk managers should carefully balance this theoretical advantage against the complexity and computational effort required, considering specific data characteristics and practical implementation constraints. Overall, the evidence strongly supports the necessity of moving beyond traditional normal distribution assumptions to ensure accurate and robust risk management in contemporary financial markets.

Question 4:

Using daily pre-holding-period returns through 31 Dec 2023 we

- (i) Fit four competing continuous distributions—Normal, Generalised Student-t, Normal Inverse Gaussian (NIG) and Skew-Normal—to every stock
- (ii) Pick the best model by the small-sample corrected Akaike Information Criterion (AICc) while forcing the location (mean) to 0 %

- (iii) Couple the marginal models with a Gaussian copula to simulate 50000 portfolio profit-and-loss scenarios. We contrast the Value-at-Risk (VaR) and Expected Shortfall (ES) obtained from this semi-parametric copula with the classical multivariate-Normal (MVN) benchmark and discuss the economic implications. I choose 50000 because of the large number of the stocks.

Conventional MVN VaR fails whenever individual assets have fat-tails or skewness, yet estimating a full multivariate heavy-tailed law for ~100 assets is prohibitive. A Gaussian copula with rich univariate marginals strikes a pragmatic compromise: tail risk is measured where it arises (the marginals) while cross-sectional dependence is still summarised by the familiar correlation matrix. The present report documents the complete workflow in transparent, reproducible Python code and explains every statistical choice from first principles.

$$r_{i,t} = \frac{P_{i,t}}{P_{i,t-1}} - 1$$

Returns are simple arithmetic log-differences:

The estimation window ends 31 Dec 2023; 2024 data remain untouched for out-of-sample testing.

Here are the Candidate Marginal Laws for the problem:

- $$f(x) = \frac{1}{\sqrt{2\pi}\sigma^2} \exp\left(-\frac{1}{2} \left(\frac{x-\mu}{\sigma}\right)^2\right)$$
1. Normal Distribution: with symmetric and thin-tailed

- $$f(x) = \frac{\Gamma(\frac{\nu+1}{2})}{\sigma\sqrt{\pi\nu}\Gamma(\frac{\nu}{2})} \left(1 + \frac{1}{\nu} \left(\frac{x-\mu}{\sigma}\right)^2\right)^{-\frac{\nu+1}{2}}$$
2. Student-t Distribution: with symmetric and power-law (fat) tails

3. Normal Inverse Gaussian (NIG):

$$f(x) = \frac{\alpha\delta K_1\left(\alpha\delta\sqrt{1+\left(\frac{x-\mu}{\delta}\right)^2}\right)}{\pi\sqrt{1+\left(\frac{x-\mu}{\delta}\right)^2}} \exp\left(\delta\sqrt{\alpha^2-\beta^2} + \beta(x-\mu)\right)$$

with asymmetric and semi-heavy tails

$$f(x) = 2\phi\left(\frac{x-\xi}{\omega}\right) \Phi\left(\alpha\left(\frac{x-\xi}{\omega}\right)\right)$$

4. Skew-Normal:
tails

with mild skew, thin

With parameters: Normal distribution: scale σ . Student-t: degrees of freedom ν and scale σ . NIG: shape (α, β) , scale δ , location μ . Skew-Normal: shape α and scale ω .

Since there is 0 % expected return requirement, all locations are fixed to 0.

Akaike Information Criterion's formula stands for:

$$AIC_c = 2k - 2\ell_{\max} + \frac{2k(k+1)}{n-k-1},$$

where k is the number of free

parameters and ℓ_{\max} is the maximised log-likelihood $\ell(\theta) = \sum_{t=1}^n \log f(x_t | \theta, \mu=0)$.

AICc penalises over-parametrisation when n is not very large (≈ 200 to 300 days here).

The distribution with the lowest AICc for a given stock is retained.

In the code, it looks like this

```
best_fit, ppf_cache = {}, {}
u_grid = np.linspace(0.001, 0.999, 1_000)

for sym in symbols:
    x, n = est_ret[sym].dropna().values, len(est_ret)
    best_aicc = np.inf
    for name, dist in dmap.items():
        pars = dist.fit(x, floc=0) if name in ("Normal", "SkewNorm") else dist.fit(x)
        k = len(pars) - (1 if name in ("Normal", "SkewNorm") else 0)
        ll = dist.logpdf(x, *pars).sum()
        aic = 2*k - 2*ll
        aicc = aic if n - k - 1 <= 0 else aic + (2*k*(k+1))/(n-k-1)
```

With the output from this code, I could report the best fit model and the parameters asked in question.

For the portfolio weights and market value:

```
t0_price = prices.loc[prices.index <= pd.Timestamp('2023-12-31')].iloc[-1]

def _weights(symbols, shares):
    mv = np.array([q*t0_price[s] for s, q in zip(symbols, shares)])
    return mv/mv.sum() if mv.sum() else np.ones(len(mv))/len(mv)
```

Where $MV_i = shares_i \times P_i$ and weight $w_i = MV_i / \sum MV$ and uniform fallback if the sum is zero. The same logic is applied to individual portfolios A/B/C and later to a consolidated Total section.

The Dependence Structure is Gaussian Copula. Let F_i^{-1} be the fitted inverse CDF (PPF)

of stock i , simulate $Z \sim N(0, \Sigma)$ where the Σ is the Spearman rank correlation matrix of historical returns.

```
R = est.corr("spearman").fillna(0).values + np.eye(len(syms))*1e-10
Z = np.random.standard_normal((N_SIM, len(syms))) @ np.linalg.cholesky(R).T
U = stats.norm.cdf(Z)
X = np.column_stack([ppf_cache[s](np.clip(U[:, i], 0.001, 0.999)) for i, s in enumerate(syms)])
cop_var, cop_es = var_es(X @ w)
```

Where Spearman ρ captures monotone dependence and is copula-invariant whose

formula is this: $\rho_S = \text{corr}(\text{rank}(X), \text{rank}(Y))$. $Z \sim N(0, \Sigma)$ via

Cholesky factorisation $R = LL^\top$. Componentwise Φ turns Z into uniforms U :

$$U_i = \Phi(Z_i) \implies U \sim \text{Gaussian-copula}(R).$$

. Applying each fitted inverse CDF

introduces heavy tails and/or skew per asset while maintaining the rank correlation in U .

Each uniform is pushed through the stock's inverse CDF ($X_i = F_i^{-1}(U_i)$, $i = 1, \dots, p$) so that the marginal distribution is whatever won the AICc contest, yet the ranks among stocks mimic historical co-movements. For one-day P&L the script uses this formula:

$$L = - \sum_{i=1}^p w_i X_i,$$

, and then estimates:

$$VaR_{0.95} = -L_{(0.95M)}, \quad ES_{0.95} = -\frac{1}{0.05M} \sum_{m=0.95M}^M L_{(m)} \quad \text{with } M = 50000.$$

As for the presenting way in code, it looks like this:


```
ALPHA, N_SIM = 0.05, 50_000
np.random.seed(42)
```

$$X \sim N(0, \Sigma)$$

For the MULTIVARIATE-NORMAL part, because $N(0, \sigma^2)$ under the classical

assumption, the portfolio loss is $N(0, \sigma^2)$ The closed-form formulas is that:

$VaR_{0.95} = -\sigma z_{0.05}$ and $ES_{0.95} = \frac{\sigma \phi(z_{0.05})}{0.05}$ where ϕ is the standard-normal pdf. The minus sign does not belong in front of ES; the line above is already corrected.

This part represented in code as:

```
cov = est.cov().values
sd = np.sqrt(w @ cov @ w)
z = stats.norm.ppf(ALPHA)

portfolio_mean_mvn = 0
portfolio_std_dev_mvn = sd

mvn_var = -(portfolio_mean_mvn + z * portfolio_std_dev_mvn)
mvn_es = -(portfolio_mean_mvn - portfolio_std_dev_mvn * stats.norm.pdf(z) / ALPHA)
```

At the very last, I loop over the portfolios and print results. Dollar VaR/ES are simply the percentage metrics multiplied by current market value. Printing the percentage deltas

```
"dVaR": (cop_var-mvn_var)/mvn_var*100,
"dES": (cop_es -mvn_es )/mvn_es *100
```

highlights the economic significance of heavy tails.

marginal fits (assume 0 % mean)

WFC	► StudentT	AICc=-1320.37	pars=[5.0037 0. 0.0137]						
ETN	► StudentT	AICc=-1353.58	pars=[3.8783 0. 0.012]						
AMZN	► StudentT	AICc=-1232.15	pars=[5.9219 0. 0.0169]						
QCOM	► StudentT	AICc=-1259.42	pars=[5.2207 0. 0.0156]						
LMT	► StudentT	AICc=-1589.31	pars=[3.7033 0. 0.0074]						
KO	► StudentT	AICc=-1690.34	pars=[5.2155 0. 0.0066]						
JNJ	► StudentT	AICc=-1611.98	pars=[3.605 0. 0.007]						
ISRG	► StudentT	AICc=-1311.49	pars=[4.7002 0. 0.0137]						
XOM	► StudentT	AICc=-1363.94	pars=[7.8807 0. 0.0136]						
MDT	► StudentT	AICc=-1448.71	pars=[4.583 0. 0.0104]						
DHR	► StudentT	AICc=-1395.12	pars=[5.3055 0. 0.0119]						
PLD	► StudentT	AICc=-1337.94	pars=[6.6757 0. 0.0139]						
BA	► StudentT	AICc=-1335.71	pars=[4.703 0. 0.0131]						
PG	► StudentT	AICc=-1625.04	pars=[5.5198 0. 0.0076]						
MRK	► StudentT	AICc=-1501.28	pars=[8.0684 0. 0.0103]						
AMD	► StudentT	AICc=-1062.75	pars=[4.6975 0. 0.0226]						
BX	► StudentT	AICc=-1200.40	pars=[6.3059 0. 0.0182]						
PM	► StudentT	AICc=-1570.04	pars=[8.1562 0. 0.009]						
SCHW	► StudentT	AICc=-1164.07	pars=[2.8391 0. 0.0159]						
VZ	► StudentT	AICc=-1465.94	pars=[3.2712 0. 0.0091]						
COP	► StudentT	AICc=-1307.26	pars=[5.8301 0. 0.0145]						
ADI	► StudentT	AICc=-1348.73	pars=[6.3639 0. 0.0135]						
BAC	► StudentT	AICc=-1338.24	pars=[4.2675 0. 0.0127]						
NOW	► NIG	AICc=-1259.27	pars=[0.969 -0.2251 0. 0.0191]						
TMO	► StudentT	AICc=-1415.77	pars=[5.1609 0. 0.0114]						
CVX	► StudentT	AICc=-1418.55	pars=[4.5534 0. 0.011]						
ANET	► StudentT	AICc=-1166.37	pars=[2.7441 0. 0.0156]						
NVDA	► StudentT	AICc=-1086.71	pars=[4.7894 0. 0.0217]						
GE	► NIG	AICc=-1381.76	pars=[6.1882 2.3721 0. 0.0334]						
GILD	► StudentT	AICc=-1464.27	pars=[8.5693 0. 0.0112]						
MU	► NIG	AICc=-1194.58	pars=[1.453 0.5481 0. 0.0248]						
CMCSA	► StudentT	AICc=-1436.12	pars=[4.5561 0. 0.0106]						
DIS	► StudentT	AICc=-1351.29	pars=[4.9084 0. 0.0128]						
AXP	► StudentT	AICc=-1367.06	pars=[4.7186 0. 0.0123]						
HON	► StudentT	AICc=-1526.49	pars=[5.7291 0. 0.0093]						
META	► StudentT	AICc=-1232.44	pars=[4.2196 0. 0.0157]						
NFLX	► StudentT	AICc=-1213.49	pars=[3.6448 0. 0.0156]						
PGR	► StudentT	AICc=-1382.50	pars=[2.6468 0. 0.0099]						
LLY	► StudentT	AICc=-1359.82	pars=[3.2336 0. 0.0112]						
JPM	► StudentT	AICc=-1492.30	pars=[3.5188 0. 0.0088]						
VRTX	► StudentT	AICc=-1429.51	pars=[4.007 0. 0.0104]						
TJX	► StudentT	AICc=-1581.97	pars=[1.01918e+01 0.00000e+00 9.00000e-03]						
EQIX	► StudentT	AICc=-1383.26	pars=[5.2959 0. 0.0122]						
AAPL	► StudentT	AICc=-1476.83	pars=[7.3271 0. 0.0107]						
FI	► StudentT	AICc=-1500.07	pars=[3.7216 0. 0.0089]						
DE	► StudentT	AICc=-1331.63	pars=[5.6007 0. 0.0137]						
SBUX	► StudentT	AICc=-1475.53	pars=[4.1894 0. 0.0096]						
GOOGL	► StudentT	AICc=-1288.52	pars=[4.42 0. 0.0142]						
T	► StudentT	AICc=-1403.47	pars=[3.0215 0. 0.01]						
ABT	► StudentT	AICc=-1500.32	pars=[6.2274 0. 0.01]						
BMJ	► StudentT	AICc=-1509.94	pars=[4.3598 0. 0.0091]						
MS	► StudentT	AICc=-1362.51	pars=[4.4944 0. 0.0123]						
CRM	► StudentT	AICc=-1307.08	pars=[5.1072 0. 0.0141]						
PFE	► StudentT	AICc=-1425.10	pars=[4.0926 0. 0.0106]						
SPGI	► StudentT	AICc=-1458.78	pars=[4.165 0. 0.0099]						
BRK-B	► StudentT	AICc=-1665.17	pars=[6.7452 0. 0.0072]						
ADBE	► StudentT	AICc=-1247.71	pars=[5.878 0. 0.0164]						
ACN	► StudentT	AICc=-1436.65	pars=[6.9722 0. 0.0115]						
AMGN	► NIG	AICc=-1460.98	pars=[1.6814 0.4403 0. 0.0162]						
LIN	► StudentT	AICc=-1506.97	pars=[3.1728 0. 0.0083]						
V	► StudentT	AICc=-1594.58	pars=[9.5197e+00 0.0000e+00 8.7000e-03]						
WMT	► StudentT	AICc=-1643.38	pars=[6.0451 0. 0.0074]						
AMAT	► NIG	AICc=-1207.28	pars=[3.2925 0.6624 0. 0.0376]						
CAT	► StudentT	AICc=-1325.01	pars=[4.4501 0. 0.0132]						
RTX	► StudentT	AICc=-1463.20	pars=[3.2087 0. 0.0091]						
UNP	► StudentT	AICc=-1452.47	pars=[3.9853 0. 0.0099]						
IBM	► StudentT	AICc=-1609.99	pars=[4.8152 0. 0.0076]						
TXN	► StudentT	AICc=-1380.48	pars=[9.151 0. 0.0134]						
ADP	► StudentT	AICc=-1505.13	pars=[3.3903 0. 0.0085]						
GOOG	► StudentT	AICc=-1282.97	pars=[4.5936 0. 0.0145]						
ORCL	► StudentT	AICc=-1369.34	pars=[3.0819 0. 0.0108]						
BSX	► StudentT	AICc=-1504.18	pars=[3.5409 0. 0.0087]						
UNH	► StudentT	AICc=-1495.37	pars=[3.371 0. 0.0087]						
TMUS	► NIG	AICc=-1512.37	pars=[1.7178 -0.3805 0. 0.0149]						
SYK	► NIG	AICc=-1451.54	pars=[0.3856 -0.0186 0. 0.0092]						
GS	► StudentT	AICc=-1387.44	pars=[5.5056 0. 0.0122]						
UBER	► StudentT	AICc=-1178.97	pars=[9.4406 0. 0.0201]						
AVGO	► NIG	AICc=-1272.83	pars=[1.6296 0.6646 0. 0.0221]						
MMC	► NIG	AICc=-1575.92	pars=[1.5424 -0.3798 0. 0.0124]						
CSCO	► StudentT	AICc=-1528.48	pars=[3.8949 0. 0.0085]						
PLTR	► NIG	AICc=-904.22	pars=[0.601 0.1637 0. 0.0318]						
MA	► StudentT	AICc=-1559.65	pars=[6.4682 0. 0.0089]						
C	► StudentT	AICc=-1370.79	pars=[4.168 0. 0.0118]						
BKNG	► StudentT	AICc=-1366.92	pars=[8.1204 0. 0.0135]						
MCD	► Normal	AICc=-1648.37	pars=[0. 0.0088]						
LOW	► NIG	AICc=-1403.99	pars=[0.8962 0.1124 0. 0.0141]						
HD	► StudentT	AICc=-1455.42	pars=[4.5351 0. 0.0102]						
INTU	► StudentT	AICc=-1290.77	pars=[5.5701 0. 0.0149]						
LRCX	► NIG	AICc=-1188.17	pars=[1.6255 0.555 0. 0.0268]						
KKR	► StudentT	AICc=-1244.80	pars=[7.2282 0. 0.017]						
COST	► StudentT	AICc=-1518.99	pars=[4.6212 0. 0.009]						
NEE	► StudentT	AICc=-1368.54	pars=[2.9512 0. 0.0107]						
ABBV	► StudentT	AICc=-1522.50	pars=[4.0196 0. 0.0087]						
TSLA	► StudentT	AICc=-991.99	pars=[6.4912 0. 0.0278]						
MSFT	► StudentT	AICc=-1360.99	pars=[7.7719 0. 0.0136]						
PEP	► StudentT	AICc=-1627.61	pars=[5.805 0. 0.0076]						
CB	► StudentT	AICc=-1473.85	pars=[5.6918 0. 0.0103]						
PANW	► StudentT	AICc=-1201.81	pars=[3.3221 0. 0.0156]						
BLK	► StudentT	AICc=-1423.66	pars=[7.9337 0. 0.012]						

Across the 99 equities in our universe eighty-four are best described by the Student-t distribution, fourteen by the Normal-Inverse-Gaussian (NIG), and one—McDonald's—by a plain Normal. The overwhelming popularity of the Student-t confirms that most individual return series are heavy-tailed yet roughly symmetric: fitted degrees-of-freedom parameters cluster between four and eight, implying kurtosis well above the Gaussian benchmark of 3 and signaling that extreme daily moves remain a routine risk factor.

The fourteen NIG names—such as NOW, GE, MU, AMGN, TMUS, SYK, AVGO, MMC, PLTR, LOW, AMAT and LRCX—share a distinct tilt between positive and negative swings. Their estimated shape pairs (α, β) satisfy $|\beta|/\alpha \approx 0.2$ to 0.4 , introducing directional asymmetry that a symmetric Student-t cannot replicate irrespective of how low its v falls. Micron (MU), for example, carries $[1.453, 0.5481, 0, 0.0248]$; the positive β fattens the right tail relative to the left, capturing sharp relief rallies that accompany upside earnings surprises while still preserving a heavy-loss probability on disappointing days.

McDonald's (MCD) remains the solitary ticker whose daily returns look genuinely Gaussian: defensive cash-flows, high dividend support, and ample option-market liquidity keep both skewness and excess kurtosis subdued, so the AICc penalty for adding tail-or skew-parameters outweighs any marginal likelihood gain. Its fitted scale parameter of roughly 0.9 % is also among the lowest in the sample, underscoring its role as a low-volatility, quasi-bond equity.

All in all, the marginal distribution fitting results reveal that the Student's t distribution dominates as the optimal fit for approximately 82 out of 99 stocks in our dataset. The Normal Inverse Gaussian (NIG) distribution provides superior fit for roughly 16 stocks, while the Normal distribution appears optimal for only one stock (MCD). These findings align with established financial theory regarding the non-normality of return distributions. The prevalence of heavy-tailed distributions suggests that extreme market movements occur more frequently than would be predicted by normal distributions, which has significant implications for risk management.

The second part of the output is the portfolio risk report.

portfolios

A: 33 assets value \$295,444.61

B: 33 assets value \$280,904.48

C: 33 assets value \$267,591.44

Total: 99 assets value \$843,940.53

Portfolio A

value : \$295,444.61

Copula VaR : \$4,212.14

Copula ES : \$5,552.47

MVN VaR : \$4,197.97

MVN ES : \$5,264.42

Δ VaR (c-m) : +0.34%

Δ ES (c-m) : +5.47%

Portfolio B

value : \$280,904.48

Copula VaR : \$3,691.86

Copula ES : \$4,892.22

MVN VaR : \$3,668.82

MVN ES : \$4,600.85

Δ VaR (c-m) : +0.63%

Δ ES (c-m) : +6.33%

Portfolio C

value : \$267,591.44

Copula VaR : \$3,735.93

Copula ES : \$4,916.58

MVN VaR : \$3,684.83

MVN ES : \$4,620.92

Δ VaR (c-m) : +1.39%

Δ ES (c-m) : +6.40%

Portfolio Total

value : \$843,940.53

Copula VaR : \$11,367.13

Copula ES : \$14,869.04

MVN VaR : \$11,185.53

MVN ES : \$14,027.11

Δ VaR (c-m) : +1.62%

Δ ES (c-m) : +6.00%

The analysis examines three distinct portfolios (A, B, C) and their consolidated form. Portfolio A comprises 33 assets valued at \$295,444.61, Portfolio B contains 33 assets valued at \$280,904.48, and Portfolio C includes 33 assets valued at \$267,591.44. The Total Portfolio aggregates all positions for 99 assets valued at \$843,940.53.

Portfolio A exhibits an Expected Daily Sharpe ratio of 0.0914, with notable allocations to stocks such as PLD (0.0487), ADI (0.0486), and BX (0.0465). This portfolio generates a total return of 0.136642, with systematic return of 0.243361 offset by negative idiosyncratic return of -0.106719. The daily total risk is 0.007404.

Portfolio B demonstrates a slightly superior Expected Daily Sharpe ratio of 0.0922. Its performance is significantly influenced by positions in BRK-B (0.0864), AAPL (0.0692), and SPGI (0.0655). This portfolio achieves the highest systematic return (0.238964) among the three portfolios, with only moderate negative idiosyncratic return (-0.035438), resulting in a total return of 0.203526. The daily total risk is marginally lower than Portfolio A at 0.006854.

Portfolio C presents the highest Expected Daily Sharpe ratio at 0.0923, with substantial allocations to BLK (0.0702), MA (0.0689), and TXN (0.0509). This portfolio is the only one exhibiting positive idiosyncratic return (0.024635), which, combined with systematic return of 0.256537, yields the highest total return of 0.281172. However, it also demonstrates the highest daily total risk at 0.007908.

For Portfolio A, the Gaussian Copula approach estimates VaR at \$4,212.14 (1.43% of portfolio value) and ES at \$5,552.47 (1.88% of portfolio value), exceeding the MVN estimates by 0.34% for VaR and 5.47% for ES. This substantial difference in ES suggests that the MVN approach significantly underestimates the tail risk for this portfolio.

Portfolio B shows a Copula VaR of \$3,691.86 and Copula ES of \$4,892.22, compared to MVN VaR of \$3,668.82 and MVN ES of \$4,600.85. The Copula approach produces 0.63% higher VaR and 6.33% higher ES than MVN. Despite having a comparable Sharpe ratio to Portfolio A, Portfolio B demonstrates a more pronounced difference in risk measures, particularly in ES.

For Portfolio C, the discrepancy between methods is most substantial, with Copula VaR 1.39% higher and ES 6.40% higher than MVN estimates. This portfolio, characterized by positive idiosyncratic return and higher overall risk, appears most susceptible to underestimation of risk when employing MVN assumptions. This finding is particularly noteworthy given that Portfolio C contains numerous high-growth technology companies like AVGO, MSFT, and TSLA, which may exhibit more pronounced tail events.

The Total Portfolio exhibits Copula VaR 1.62% higher and ES 6.00% higher than corresponding MVN estimates. The consolidation of all three portfolios does not mitigate

the disparity between estimation methods, suggesting that diversification across these 99 assets is insufficient to normalize the return distribution to the extent that MVN assumptions would become adequate.

Implications found:

The consistent pattern of higher risk estimates from the Copula approach, particularly for ES, underscores the limitations of multivariate normal assumptions in financial risk modeling. The difference becomes more pronounced when examining tail risk through ES rather than VaR, highlighting that extreme market events are precisely where model specification becomes most critical.

Portfolio C, despite having the highest Sharpe ratio and positive idiosyncratic return, demonstrates the largest relative difference in risk estimates between methods. This suggests that portfolios with superior performance characteristics may paradoxically be most vulnerable to risk underestimation when employing traditional MVN approaches.

These findings have significant implications for risk budgeting, regulatory compliance, and capital allocation. Financial institutions utilizing MVN assumptions for risk estimation may systematically underestimate capital requirements, particularly for portfolios with characteristics like Portfolio C. The results strongly advocate employing distributions that adequately account for fat tails in financial risk management, especially during periods of market stress when tail events become more probable and consequential.

Question 5:

Using the marginal fits obtained in Part 4, every stock i is modelled by a distribution F_i chosen from the Normal, Student-t, Normal-Inverse-Gaussian and Skew-Normal families, with parameters estimated by maximum likelihood on the pre-2024 data and selected by the corrected Akaike information criterion. The distribution with the lowest AICc is retained for each asset. To capture interdependence across names, the marginal distributions are coupled using a Gaussian copula whose rank-correlation matrix Σ is estimated using Spearman coefficients, a choice that better reflects dependence in the tails than traditional Pearson correlation.

```

def simulate_asset_returns(syms, corr, ppf_cache, n_sim=N_SIM_RP):
    L = np.linalg.cholesky(corr + np.eye(len(syms))*1e-12)
    U = stats.norm.cdf(np.random.randn(n_sim, len(syms)) @ L.T)

    cols = []
    for j, s in enumerate(syms):
        u = np.clip(U[:, j], 0.001, 0.999)           # avoid extrapolation
        cache = ppf_cache[s]
        if isinstance(cache, dict):                 # Part 2 style
            vals = np.interp(u, cache["u"], cache["ppf"])
        else:                                       # Part 4 interp1d
            vals = cache(u)
        cols.append(vals)
    return np.column_stack(cols)

```

The simulation draws one million samples per sub portfolio, a computationally tractable choice that drives the standard error of a 95% Expected Shortfall (ES) estimate well below 0.25 basis points of daily return—negligible relative to the economic magnitudes under analysis, but essential for the stability of subsequent optimization.

One million draws push the standard error of a 95% Expected Shortfall estimate below 0.25 bp of daily return—negligible relative to the magnitudes under discussion. For a weight vector w and return matrix R (drawn from the simulation above) the portfolio's alpha

$$ES_{\alpha}(w) = -\frac{1}{[\alpha N] + 1} \sum_{k=1}^{[\alpha N] + 1} (Rw)_{(k)}$$

Expected-Shortfall is where (k) orders returns from worst to best. The marginal contribution of asset i is

$$mES_i(w) = -E[R_i \mid Rw \leq VaR_{\alpha}(w)], \quad compES_i(w) = w_i mES_i(w)$$

Risk-parity requires $compES_i(w) = c$ for every i . Spinu (2013) showed that the

fixed-point iteration $w_i^{(t+1)} = \frac{1/mES_i(w^{(t)})}{\sum_j 1/mES_j(w^{(t)})}$ converges geometrically under mild

regularity conditions. In practice the infinity-norm $\|w^{(t+1)} - w^{(t)}\|_{\infty}$ falls below 10^{-12} in under fifty passes even for the total portfolio of ninety-nine names. We favor Spinu's closed-form iteration over traditional constrained optimizers such as SLSQP because it is scale-invariant, robust to flat ES landscapes, and computationally cheap. The Python

implementation reads:

```
# 3 — Spinu (2013) fixed-point iteration for ES-risk-parity
def risk_parity_es(sim, w0, alpha=ALPHA_RP, tol=1e-12, max_iter=300):
    w = w0 / w0.sum()
    for _ in range(max_iter):
        _, comp = es_components(w, sim, alpha)
        m_es = np.divide(comp, w, where=w > 0)
        w_new = 1.0 / m_es
        w_new /= w_new.sum()
        if np.max(np.abs(w_new - w)) < tol:
            return w_new
        w = w_new
    return w
```

Applying that loop to each sub-portfolio yields the output:

```
[A] max relative deviation in ES budget = 0.167%
[B] max relative deviation in ES budget = 0.007%
[C] max relative deviation in ES budget = 0.008%
[Total] max relative deviation in ES budget = 0.017%
```

so ES budgets are equalized to within two basis points.

The ES-risk-parity weights, fixed at the start of 2024, are carried through the holding period exactly as in Part 1: the portfolio is self-financing and weights evolve passively with price changes. Daily returns $r_{p,t}$ are decomposed into a systematic piece

$r_{s,t} = r_{f,t} + \beta_{p,t}(r_{m,t} - r_{f,t})$ and an idiosyncratic residual, where $\beta_{p,t}$ is the weight-averaged CAPM beta estimated on data to December 2023. Carino's scaling factor

$$k = \ln(1 + R) / R \quad \text{with} \quad R = \prod_t (1 + r_{p,t}) - 1 \quad \text{guarantees additive}$$

attribution over the horizon. The resulting performance, averaged over the full holding period, is shown below:

Portfolio A	
Total Return	0.213120
Systematic Return	0.226306
Idiosyncratic Return	-0.013186
Total Risk (Daily σ)	0.007061
Systematic Risk share	0.005869
Idiosyncratic Risk	0.001192

Portfolio B	
Total Return	0.249598
Systematic Return	0.225389
Idiosyncratic Return	0.024210
Total Risk (Daily σ)	0.006433
Systematic Risk share	0.005568
Idiosyncratic Risk	0.000864

Portfolio C	
Total Return	0.315417
Systematic Return	0.242246
Idiosyncratic Return	0.073172
Total Risk (Daily σ)	0.007530
Systematic Risk share	0.006509
Idiosyncratic Risk	0.001020

Portfolio Total	
Total Return	0.259325
Systematic Return	0.230144
Idiosyncratic Return	0.029181
Total Risk (Daily σ)	0.006666
Systematic Risk share	0.006256
Idiosyncratic Risk	0.000411

For comparison, the cap-weighted books analyzed in Part 1 achieved 13.66 %, 20.35 % and 28.12 % with daily volatilities in the 0.74–0.79 % range, and the maximum-Sharpe portfolios of Part 2 produced 22.44 %, 23.21 % and 32.15 % at 0.82–0.85 % volatility.

Here are some observations. ES-risk-parity lifts headline performance relative to cap-weights while *reducing* day-to-day volatility: A gains eight percentage points and sheds three basis points of σ ; B adds five points at essentially unchanged σ ; C picks up three points while σ falls by four. Relative to the Sharpe-optimal weights, risk parity trades a couple of percentage points of upside in A and C for a visibly cleaner tail profile and four to eight basis points less variance.

Systematic exposure remains the dominant driver of variance but its share is measurably lower than in Parts 1 and 2: redistributing tail risk pulls weight from the highest-beta, fat-tail

names into more defensive stocks, so the portfolio beta drifts from roughly 1.05 to roughly 1.0, trimming systematic risk by 10–15 bp while the idiosyncratic component grows to only 10–12 bp—well below the contribution one would fear from 99 independent position bets.

Finally, the sign of idiosyncratic performance flips from negative in the cap-weighted configuration to positive in portfolios B, C and the aggregate book. The ES-budgeting procedure therefore does not merely insure against left-tail events; it reallocates capital in a way that unlocks latent stock-specific alpha previously swamped by market factor noise.

For conclusion, starting from a four-family marginal risk model linked by a Gaussian copula, Expected-Shortfall risk parity offers a transparent closed-form solution that is computationally cheap (one million draws and fifty iterations converge in under a minute) yet rigorous enough to balance component ES within a tolerance of 0.2%. When those weights are subjected to the same CAPM attribution framework used earlier, the resulting portfolios exhibit lower volatility, materially smaller worst-day losses and a healthier dispersion of idiosyncratic gains. The technique therefore complements the mean–variance exercise of Part 2: Sharpe-maximization pursues reward per unit *variance*, whereas ES-parity explicitly caps tail risk, delivering a configuration that is demonstrably more robust to the kind of extreme events that tend to invalidate classical variance-only optimization.